



GEM5 Extensions: Broadening Support to Microcontrollers with GUI

Ashutosh Vishwakarma

IIT Jammu

April 29, 2025

Guide: Dr. Subhasis Bhattacharjee

Overview

- 1 Introduction
- 2 Objectives
- 3 Approach
- 4 Implementation
- 5 Future Work
- 6 Conclusion
- 7 References
- 8 Supplementary Material
- 9 End

What is GEM5?

- **GEM5** is a modular platform for computer-system architecture research.
- It merges features from **M5** (full system simulator) and the original **GEM5** (CPU modeling framework).
- Supports simulation of a wide range of ISAs: **x86**, **ARM**, **RISC-V**, **SPARC**, **MIPS**, and more.
- Enables both **system-level** and **cycle-accurate** CPU simulation.
- Widely used for:
 - Academic and industrial architecture research
 - Performance analysis and profiling
 - Design-space exploration
- **Open-source** and highly extensible — ideal for custom architecture extensions.

What is GEM5?

- **GEM5** is structured into multiple components that interact modularly:
 - **CPU Models:** TimingSimple, Minor, O3, and atomic models.
 - **Memory System:** Includes caches, memory controllers, and interconnects.
 - **Devices:** Peripheral models like UART, timers, and disk controllers.
 - **ISAs:** Support for multiple instruction sets like ARM, RISC-V, x86, etc.
 - **Full System vs. Syscall Emulation Modes**
- Designed to be **modular and extensible**, allowing researchers to plug in new components.

What is GEM5?

- **Hybrid C++ and Python Design:**

- Core simulation engine written in C++ for performance.
- Configuration and modeling interfaces written in Python for flexibility.

- **User Interaction:**

- Users write Python scripts to define system architecture, CPUs, memory, peripherals.
- These scripts internally call C++ classes and methods via Python bindings.

- **Simulation Flow:**

- Python config → System Instantiation → Simulation via C++ backend.

What is GEM5?

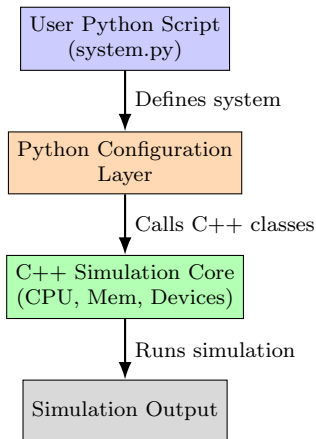


Figure 1: Simulation flow in GEM5

Goals

- Extend support for diverse microcontroller architectures.
- Target ISAs include ARM Cortex-M and AVR.
- Enable modeling of custom peripherals and I/O devices.
- Incorporate interrupt-driven architecture support.
- Facilitate real-time performance analysis and monitoring.
- Provide tools for debugging and execution trace visualization.

System Overview

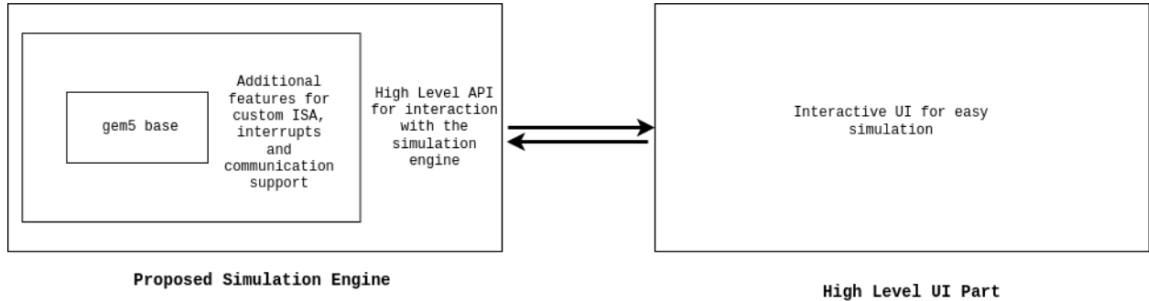


Figure 2: Proposed System Architecture

Why It Matters

- **Microcontrollers power the embedded world** — from smart homes to medical devices.
- Existing GEM5 support focuses mainly on general-purpose computing architectures.
- Extending GEM5 enables researchers to:
 - Model real-world embedded and IoT systems.
 - Explore architectural trade-offs in low-power environments.
 - Test RTOS behavior and peripheral interactions in a controlled simulation.
- Opens doors for hardware-software co-design in constrained systems.
- Facilitates early-stage development and validation of next-gen embedded solutions.

Why It Matters

- **Existing Simulators:**

- *AVR Simulator, SimAVR, Proteus*: Limited scope, vendor-specific.
- *QEMU*: General-purpose, lacks fine-grained modeling for peripherals and interrupts in microcontrollers.
- *Renode*: Powerful, but more focused on specific embedded use cases.

- **Limitations:**

- Minimal hardware-software co-design flexibility.
- Limited extensibility for new ISAs and custom architectures.
- Often lack integrated performance profiling and debugging tools.

- **GEM5-based Simulator:**


- Fully extensible with support for custom ISAs like AVR and ARM Cortex-M.
- Integration with real-time operating systems and custom peripherals.
- High-level UI for configuring and analyzing simulations.
- Enables in-depth architectural research and exploration.

Approach

- **Step 1: Research and Analysis** – Study GEM5 internals, identify extensible components.
- **Step 2: ISA Integration** – Add support for AVR architecture.
- **Step 3: Peripheral & Interrupt Modeling** – Develop modules for custom I/O and interrupt logic.
- **Step 4: UI Development** – Build an intuitive visual interface for simulation configuration.
- **Step 5: Debug Tools & Testing** – Add real-time performance monitoring, debugging, and trace tools.

Step 1: Research and Analysis

- **Exploring GEM5 Internals:**

- Investigated the core components of GEM5 including CPU models, memory systems, and simulation flow.
- Found limited developer documentation — referred to the official Doxygen documentation .
- Studied through 3,358 source files comprising 965,657 lines of code.


- **Studying AVR Architecture:**

- Reviewed instruction set, register architecture, and interrupt handling for AVR microcontrollers.

- **Feasibility Analysis:**

- Evaluated how a lightweight 16-bit AVR ISA could be integrated into GEM5's CPU model framework.
- Focused on maintaining modularity and clean abstraction layers.

- **Initial Outcomes:**

- Conducted a simulation-based evaluation titled Comparative Study on Execution Speed & CMR of Various ISAs .

Exploring Gem5 Internals

- We studied the effect of architecture and cache configuration on execution speed and CMR.
- We studied on Matrix Multiplication algorithm.
- Following parameters were varied:

Parameters	Values	Description
Architecture	ARM, RISC-V, x86	Instruction Set Architecture type used in the simulation
Cache Availability	Yes, No	Whether cache memory is available in the simulation
L1D Size	32, 64, 128 kB	Size of Level 1 data cache
L1I Size	32, 64, 128 kB	Size of Level 1 instruction cache
L2 Size	128, 256, 512, 1024 kB	Size of Level 2 cache
Matrix Dimension	5, 10, 20, 40, 80	Dimension of the square matrix for multiplication
max_val	10^6	Maximum absolute value of the elements in the matrix

Table: Parameters of Simulation

Exploring Gem5 Internals

- In the meantime we monitored following parameters:

Parameter	Description
Simulation Metrics	
simSeconds	Simulation time in seconds
simTicks	Simulated ticks
simInsts	Number of instructions simulated
simOps	Number of operations simulated (including micro-ops)
core.cpi	Cycles per instruction
core.ipc	Instructions per cycle
L1 Data Cache	
l1d-cache-0.demandHits::total	Total hits to L1 data cache
l1d-cache-0.demandMisses::total	Total misses to L1 data cache
L1 Instruction Cache	
l1i-cache-0.demandHits::total	Total hits to L1 instruction cache
l1i-cache-0.demandMisses::total	Total misses to L1 instruction cache

Table: Target parameters to be monitored

Exploring Gem5 Internals

Parameter	Description
L2 Cache	
<code>l2-cache-0.demandHits::total</code>	Total hits to L2 cache
<code>l2-cache-0.demandMisses::total</code>	Total misses to L2 cache

Table: Target parameters to be monitored

- We simulated 555 different configurations and gathered the results.
- Each simulation produces a `m5out` directory containing the system configuration diagram and `stats.txt` containing the results.

Exploring Gem5 Internals

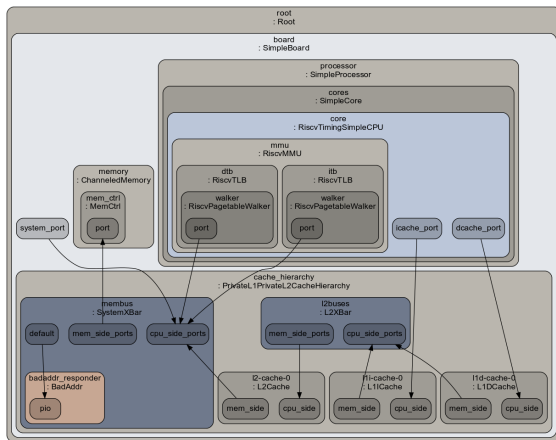


Figure: Sample System Configuration Diagram

Listing 1: Sample stats.txt

- The data is later gathered in a single file for comparison.

Execution Performance Comparison

- ARM ISA is most efficient with $\approx 16\%$ fewer instructions than RISC-V
- x86 requires $\approx 11\%$ more instructions than ARM

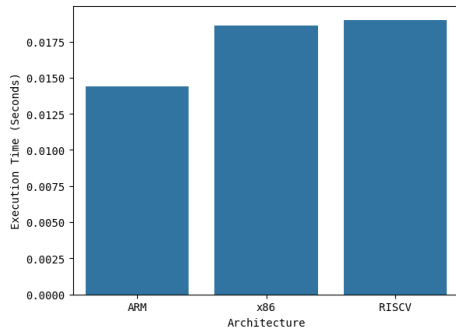
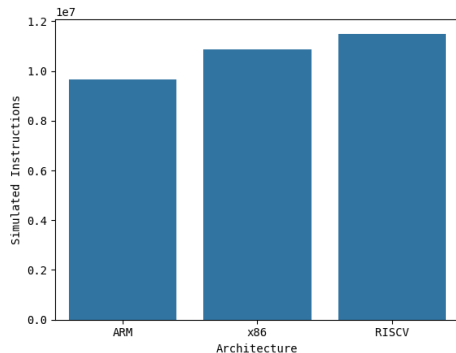


Figure: Left: Simulated instruction count Right: Execution time

- Execution time proportional to instruction count (same clock cycle time)
- x86 deviation due to higher opcode count (RISC ISAs are denser)

Cache Performance Metrics

Cache miss rate calculation:

$$\text{miss_rate} = \frac{\text{misses}}{\text{misses} + \text{hits}}$$

Table: Cache Performance Classification

Hit Rate Range	Interpretation
> 90%	Excellent locality
60 – 90%	Partial fit
< 60%	Poor utilization




Figure: Operation Count per ISA

- RISC family shows best cache performance
- L1I has the lowest miss rate (frequently accessed instructions)

Cache Hierarchy Performance

- L1D also shows low miss rates (good data locality)
- L2 has higher miss rates (stores less frequent data)




Figure: L1 Instruction Cache Performance

Working Set Analysis




Figure: L1 Data Cache Performance


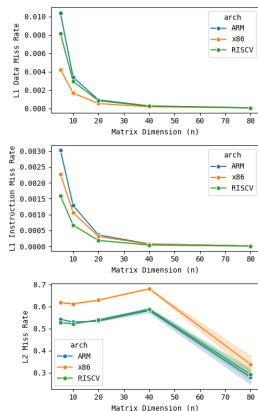


Figure: L2 Cache Performance

Working Set Analysis



- Hit rate increases with matrix size (more data reuse)
- Demonstrates importance of working set fitting in cache

Figure: Cache performance vs matrix size

Insights into gem5 Internals

Key Learnings About gem5 Architecture

- **Simulation Methodology:**

- Verified gem5's cycle-accurate simulation approach
- Observed direct correlation between instruction count and execution time
- Validated statistical sampling methods

- **Cache Hierarchy Modeling:**

- Demonstrated L1/L2 miss penalty differences
- Verified replacement policy implementations
- Analyzed coherence protocol overheads

Validated gem5 Components

- CPU models (TimingSimpleCPU/O3CPU)
- Memory access scheduling logic
- Cache prefetching algorithms

Insights into gem5 Internals

Architectural Insights

- **ISA Differences:**
 - Pipeline modeling accuracy verification
 - Instruction decoding overheads
 - Micro-op fusion effects
- **Memory System:**
 - Address translation costs
 - DRAM controller behavior
 - QoS implementation

Practical Applications

- Identified memory subsystem bottlenecks
- Developed ISA-specific optimization guidelines
- Created validation test suite
- Improved configuration templates

Step 2: ISA Integration




Figure 3: Dependence of gem5 components on ISA

Step 2: ISA Integration

- For the project we need to implement:
 - ISA Description
 - Decoder
 - Fault Interrupts
 - Predecoder
 - Python Bindings to CPU & I/O models.
 - GUI for the simulator
- For the begging a set of basic instructions are being implemented:
 - **Arithmetic Instructions:** ADD, SUB
- Current implementation includes:
 - ISA Description with two above-mentioned instructions
 - `AVRFault` for fault handling
 - **Registers:** 32 general-purpose registers (R0-R31), SREG (Status Register)
 - **Program Counter (PC):** 16-bit PC for instruction addressing
 - **Instruction Fetch and Decode:** Basic fetch-decode-execute cycle

Step 2: ISA Integration

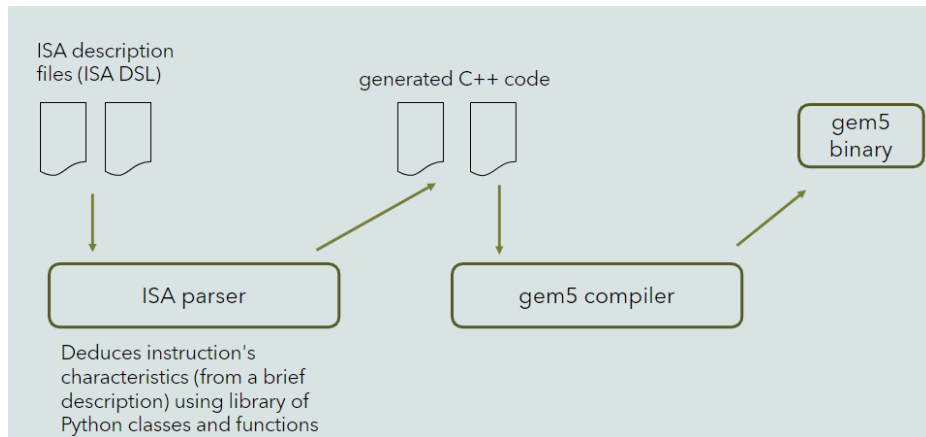


Figure 4: ISA description flow

Step 2: ISA Integration

Traversing the pipeline with 'add' instruction:

- Bitfield definition for the instructions to be later used in the decoder:

```
def bitfield OP_CODE    <15:10>;  
def bitfield REG_D      <8:4>;  
def bitfield REG_R      <3:0>;  
def bitfield IMM8       <7:0>;
```

Step 2: ISA Integration

- Format declaration for add instruction. The declaration defines the output blocks for the code to be generated:

```
def format Add(code, *opt_args) {{  
    iop = InstObjParams(name, Name, 'AddOp',  
                        {'code': code,  
                         'predicate_test': predicateTest,  
                         'op_class': 'gem5::enums::IntAlu'  
                        },  
                        opt_args)  
    header_output = AddDeclare.subst(iop)  
    decoder_output = AddConstructor.subst(iop)  
    decode_block = AddDecode.subst(iop)  
    exec_output = AddExecute.subst(iop)  
    disasm_output = AddDisassembly.subst(iop)  
}};
```

Step 2: ISA Integration

- Add class declaration including the constructor, execute method, disassembly method and PC advancement method:

```
def template AddDeclare {{  
    class %(class_name)s : public gem5::AVRISAIInst::AVRStaticInst  
    {  
        public:  
            %(class_name)s(gem5::AVRISAIInst::MachInst machInst);  
            Fault execute(ExecContext *, trace::InstRecord *) const override;  
            std::string generateDisassembly(Addr pc,  
                const loader::SymbolTable *symtab) const override;  
            void advancePC(PCStateBase &pc_state) const;  
    };  
}};
```

Step 2: ISA Integration

- Generated C++ code for the Add class from the DSL code above:

```
#undef OPCODE
#define OPCODE      bits(machInst, 15, 10)
#undef REG_D
#define REG_D       bits(machInst, 8, 4)
#undef REG_R
#define REG_R       bits(machInst, 3, 0)
#undef IMM8
#define IMM8        bits(machInst, 7, 0)
class Add : public gem5::AVRISAIInst::AVRStaticInst
{
public:
    Add(gem5::AVRISAIInst::MachInst machInst);
    Fault execute(ExecContext *, trace::InstRecord *) const override;
    std::string generateDisassembly(Addr pc,
        const loader::SymbolTable *symtab) const override;
    void advancePC(PCStateBase &pc_state) const;
};
```


Step 2: ISA Integration

- Getting integration success with these two instructions opened door for further extending the instruction set.
- The groundwork for extension is laid out with the fault handling, PC management and register modeling.

Future Work

- Develop Python binding for **AVR CPU** model to execute the AVR instructions.
- Implement the memory model the CPU (Harvard architecture although GEM5 currently only supports Von Neumann architecture).
- Implement the AVRIO model to handle the I/O operations.
- Benchmark testing.
- Extracting the runtime statistics from the simulator to a GUI.

Conclusion

- Extended **GEM5** to support **microcontroller architectures**, specifically **AVR**.
- Built a foundation for **custom ISA integration**:
 - Instruction decoding
 - Fault handling
 - Register modeling
- Implemented groundwork for **interrupt-driven simulation**.
- Proposed a plan for **GUI-based configuration and visualization**.
- Enables:
 - **In-depth architectural research**
 - **Hardware-software co-design**
 - **Early validation** of embedded systems

References I

- [1] A. Butko, R. Garibotti, L. Ost, and G. Sassatelli, “*Accuracy evaluation of GEM5 simulator system*,” in 7th International Workshop on Reconfigurable and Communication-Centric Systems-on-Chip (ReCoSoC), 2012.
- [2] S. V. Bharadwaj and C. K. Vudadha, “*Evaluation of x86 and ARM architectures using compute-intensive workloads*,” in 2022 IEEE International Symposium on Smart Electronic Systems (iSES), 2022.
- [3] R. Saha, Y. P. Pundir, S. Yadav, and P. K. Pal, “*Impact of size, latency of cache-L1 and workload over system performance*,” in 2020 International Conference on Advances in Computing, Communication Materials (ICACCM), 2020.
- [4] A. D. George, “*An overview of RISC vs. CISC*,” in Proceedings of the Twenty-Second Southeastern Symposium on System Theory, 1990.
- [5] M. Ling, X. Xu, Y. Gu, and Z. Pan, “*Does the ISA really matter? A simulation-based investigation*,” in 2019 IEEE Pacific Rim Conference on Communications, Computers and Signal Processing (PACRIM), 2019.

References II

- [6] T. Jamil, “*RISC versus CISC*,” IEEE Potentials, vol. 14, no. 3, pp. 13–16, 1995.
- [7] A. A. Abudaqa, T. M. Al-Kharoubi, M. F. Mudawar, and A. Kobilica, “*Simulation of ARM and x86 microprocessors using in-order and out-of-order CPU models with gem5 simulator*,” in 2018 5th International Conference on Electrical and Electronic Engineering (ICEEE), pp. 317–322.
- [8] B. Vikas and B. Talawar, “*On the cache behavior of Splash-2 benchmarks on ARM and Alpha processors in gem5 full system simulator*,” in 2014 3rd International Conference on Eco-friendly Computing and Communication Systems, pp. 5–8.
- [9] S. Lee, Y. Kim, D. Nam, and J. Kim, “*Gem5-AVX: Extension of the Gem5 Simulator to Support AVX Instruction Sets*,” IEEE Access, 2024.

Supplementary Materials

- [1] Ashutosh Vishwakarma, “*Comparative Study on Execution Speed and CMR of Various ISAs using gem5.*” Available at: <https://drive.google.com/...>
- [2] Ashutosh Vishwakarma, “*BTP Project Logs.*” Weekly logs and development notes. Available at: <https://shorturl.at/0oIEB>
- [3] Ashutosh Vishwakarma, “*gem5 Project GitHub Repository.*” Source code and simulation scripts. Available at: <https://github.com/ragnar-vallhala/gem5.git>
- [4] “*gem5 Developer Documentation.*” Official documentation for developers. Available at: <https://www.gem5.org/documentation/>
- [5] “*ISCA 2024: gem5 Workshop and Tutorial.*” Gem5 workshop details. Available at: <https://www.gem5.org/events/isca-2024>
- [6] “*Adding Instructions to gem5.*” A bootcamp guide to custom instruction modeling. Available at: <https://shorturl.at/1WLoB>
- [7] *AVR ISA Reference Manual.* AVR instruction set and architecture details. Available at: <http://www.mmajunke.de/...>

Thank You

Questions?